

# Knn on amazon reviews

June 6, 2018

```
In [0]: import sqlite3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.plotly as py
import plotly.graph_objs as go
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.decomposition import TruncatedSVD
```

```
In [0]: final = pd.read_csv("final.csv")
```

## 1 Sampling the data

```
In [0]: positive_reviews = final.loc[final["Score"] == "positive"]
negative_reviews = final.loc[final["Score"] == "negative"]

positive_reviews = positive_reviews.sample(n = 5000)
negative_reviews = negative_reviews.sample(n = 5000)

f1 = [positive_reviews, negative_reviews]
final_data = pd.concat(f1)

final_data = final_data.drop(["Text"], axis = 1)
final_data = final_data.drop(final_data.columns[0], axis = 1)
```

## 2 Sorting the dataset based on 'time'

```
In [0]: labels = final_data.Score
In [0]: final_data = final_data.sort_values("Time")
In [49]: final_data.shape
Out[49]: (10000, 10)
```

### 2.1 Splitting the dataset into train and test data

```
In [0]: n = final_data.shape[0]
        train_size = 0.7

        train_set = final_data.iloc[:int(n*train_size)]
        test_set = final_data.iloc[int(n*train_size):]

        X_train = train_set.CleanedText
        y_train = train_set.Score

        X_test = test_set.CleanedText
        y_test = test_set.Score

In [11]: import re
        # Tutorial about Python regular expressions: https://pymotw.com/2/re/
        import string
        from nltk.corpus import stopwords
        from nltk.stem import PorterStemmer
        from nltk.stem.wordnet import WordNetLemmatizer

        stop = set(stopwords.words('english')) #set of stopwords
        sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

        def cleanhtml(sentence): #function to clean the word of any html-tags
            cleanr = re.compile('<.*?>')
            cleantext = re.sub(cleanr, ' ', sentence)
            return cleantext
        def cleanpunc(sentence): #function to clean the word of any punctuation or special char
            cleaned = re.sub(r'[?|!|\'|\"|#]',r'',sentence)
            cleaned = re.sub(r'[.,|)|(|\\|/]',r'',cleaned)
            return cleaned
        print(stop)
        print('*****')
        print(sno.stem('tasty'))

{"don't", 'this', 'me', 'if', 'hers', 'at', 'should', 'from', 'now', 'myself', 'or', 'do', 'ab
*****
tasti
```

### 3 K-NN using 'Bag of Words'

```
In [0]: count_vect = CountVectorizer() #in scikit-learn
X_train1 = count_vect.fit_transform(X_train)
X_test1 = count_vect.transform(X_test)
```

```
In [13]: #Standardization
from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train1 = sc.fit_transform(X_train1)
X_test1 = sc.transform(X_test1)
```

/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:475: DataConversionWarning:  
Data with input dtype int64 was converted to float64 by StandardScaler.

```
In [14]: from sklearn.model_selection import TimeSeriesSplit
# creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train1)]
# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train1, y_train, cv = my_cv, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)


# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()
```

```
print("the misclassification error for each k value is : ", (np.round(MSE,3)))
```

The optimal number of neighbors is 21.



the misclassification error for each k value is : [0.452 0.441 0.439 0.426 0.426 0.425 0.406 0.392 0.395 0.404 0.417 0.436 0.452 0.458 0.467 0.472 0.477 0.479 0.482 0.484]

### 3.0.1 Best 'k' = 21

```
In [0]: k = 21
knn = KNeighborsClassifier(n_neighbors = k)
knn.fit(X_train1,y_train)
pred = knn.predict(X_test1)
```

## 3.1 Generate the confusion matrix

```
In [22]: import itertools
from sklearn.metrics import confusion_matrix
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
x = knn.predict(X_train1)
```

```

tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)
print('\n****Train accuracy for k = {} is {:.2f}'.format(k,tr_acc))
print('\n****Test accuracy for k = {} is {:.2f}'.format(k,acc))
print(confusion_matrix(y_test, pred))

```

```
****Train accuracy for k = 21 is 77.46
```

```
****Test accuracy for k = 21 is 64.67
```

```
[[991 622]
 [438 949]]
```

## 4 K-NN using Tf-idf

In [0]: *#TF-IDF*

```

tf_idf_vect = TfidfVectorizer(ngram_range=(1,2))
X_train2 = tf_idf_vect.fit_transform(X_train)
X_test2 = tf_idf_vect.transform(X_test)

```

In [0]: *#Standardization*

```

from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train2 = sc.fit_transform(X_train2)
X_test2 = sc.transform(X_test2)

```

In [84]: *# creating odd list of K for KNN*

```

myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

```

*# empty list that will hold cv scores*

```
cv_scores = []
```

```
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train2)]
```

*# perform 10-fold cross validation*

```
for k in neighbors:
```

```
    knn = KNeighborsClassifier(n_neighbors=k)
```

```
    scores = cross_val_score(knn, X_train2, y_train, cv = my_cv, scoring='accuracy')
```

```
    cv_scores.append(scores.mean())
```

*# changing to misclassification error*

```
MSE = [1 - x for x in cv_scores]
```

*# determining best k*

```
optimal_k = neighbors[MSE.index(min(MSE))]
```

```
print('\nThe optimal number of neighbors is %d.' % optimal_k)
```

*# plot misclassification error vs k*

```

plt.plot(neighbors, MSE)


for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", (np.round(MSE,3)))

```

The optimal number of neighbors is 19.



```

the misclassification error for each k value is : [0.507 0.494 0.503 0.472 0.469 0.48 0.488
0.48 0.491 0.492 0.492 0.492 0.492 0.492 0.492 0.492 0.492 0.492 0.492 0.492
0.492]

```

#### 4.0.1 Best 'k' = 19

```

In [85]: k = 19
         knn = KNeighborsClassifier(n_neighbors = k)
         knn.fit(X_train2,y_train)

```

```

pred = knn.predict(X_test2)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
x = knn.predict(X_train2)
tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)
print('\n****Train accuracy for k = {} is {:.2f}'.format(k,tr_acc))
print('\n****Test accuracy for k = {} is {:.2f}'.format(k,acc))

```

\*\*\*\*Train accuracy for k = 19 is 48.40

\*\*\*\*Test accuracy for k = 19 is 53.97

#### 4.1 Generate the confusion matrix

```
In [0]: print(confusion_matrix(y_test, pred))
```

```

[[ 0 1608]
 [ 0 1392]]

```

### 5 Preparing Data using Word2Vec for K-NN

```
In [0]: # Train your own Word2Vec model using your own text corpus
```

```

import gensim
i=0
train_sent=[]
for sent in X_train:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())
            else:
                continue
    train_sent.append(filtered_sentence)

```

```

In [0]: i=0
test_sent=[]
for sent in X_test:
    filtered_sentence=[]
    sent=cleanhtml(sent)
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if(cleaned_words.isalpha()):
                filtered_sentence.append(cleaned_words.lower())

```

```

        else:
            continue
    test_sent.append(filtered_sentence)

```

```

In [0]: from gensim.models import Word2Vec
        from gensim.models import KeyedVectors

```

```

w2v_model1=gensim.models.Word2Vec(train_sent,min_count=5,size=50, workers=4)
w2v_model2=gensim.models.Word2Vec(test_sent,min_count=5,size=50, workers=4)

```

## 5.1 K-NN using Average Word2Vec

```

In [54]: #AVG-W2V

```

```

sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sent in train_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model1.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

7000

50

```

In [55]: #AVG-W2V

```

```

sent_vectors2 = []; # the avg-w2v for each sentence/review is stored in this list
for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model2.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    sent_vec /= cnt_words
    sent_vectors2.append(sent_vec)
print(len(sent_vectors2))
print(len(sent_vectors2[0]))

```



3000  
50

```
In [0]: #Standardization
from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train3 = sc.fit_transform(sent_vectors)
X_test3 = sc.transform(sent_vectors2)

In [57]: # creating odd list of K for KNN
myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train3)]
# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train3, y_train, cv = my_cv, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

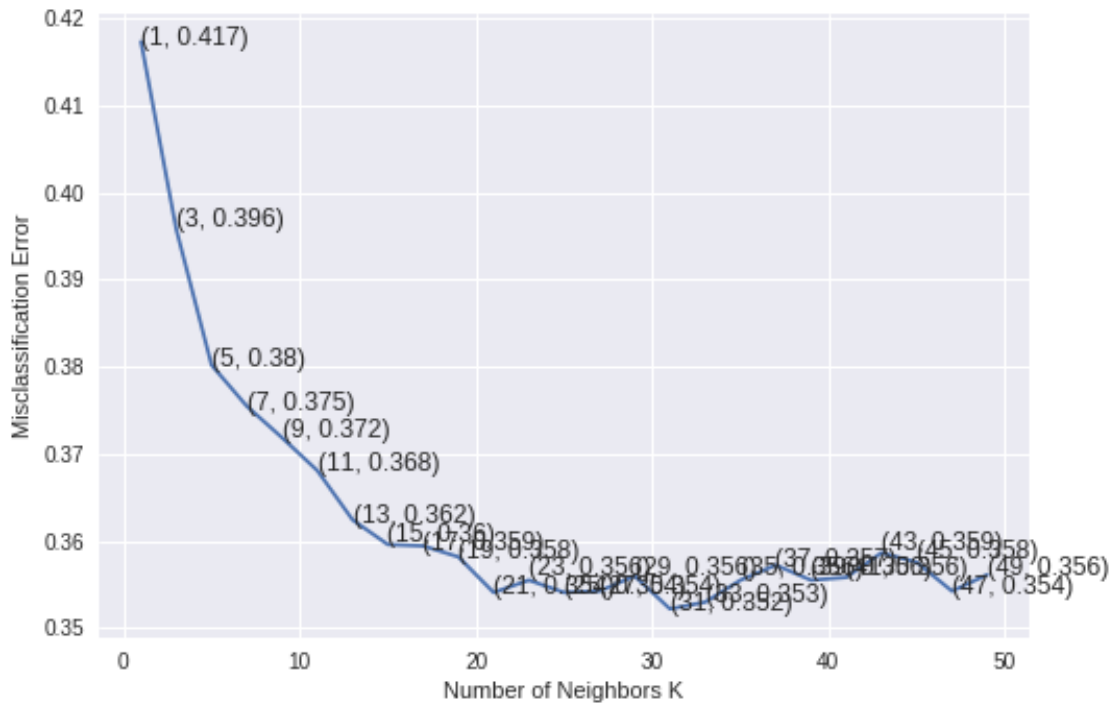
# plot misclassification error vs k
plt.plot(neighbors, MSE)

for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", (np.round(MSE,3)))
```

The optimal number of neighbors is 31.



the misclassification error for each k value is : [0.417 0.396 0.38 0.375 0.372 0.368 0.362 0.354 0.354 0.356 0.352 0.353 0.356 0.357 0.356 0.356 0.359 0.358 0.354 0.356]

### 5.1.1 Best 'k' = 31

```
In [58]: k = 31
         knn = KNeighborsClassifier(n_neighbors = k)
         knn.fit(X_train3,y_train)
         pred = knn.predict(X_test3)
         acc = accuracy_score(y_test, pred, normalize=True) * float(100)
         x = knn.predict(X_train3)
         tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)
         print('\n****Train accuracy for k = {} is {:.2f}'.format(k,tr_acc))
         print('\n****Test accuracy for k = {} is {:.2f}'.format(k,acc))
```

```
****Train accuracy for k = 31 is 68.73
```

```
****Test accuracy for k = 31 is 51.50
```

## 5.2 Generate the confusion matrix

```
In [59]: print(confusion_matrix(y_test, pred))
```

```
[[1428 133]
 [1322 117]]
```

## 6 K-NN using Word2Vec-Tfidf

In [0]: *#TF-IDF*

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf = tf_idf_vect.fit_transform(X_train)
tfidf_feat = tf_idf_vect.get_feature_names()
    # tfidf words/col-names
    # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in train_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model1.wv[word]
            # obtain the tf_idf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf

        except:
            pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    train_vectors.append(sent_vec)
    row += 1
```

In [0]: *#TF-IDF*

```
tf_idf_vect = TfidfVectorizer()
final_tf_idf = tf_idf_vect.fit_transform(X_test)
tfidf_feat = tf_idf_vect.get_feature_names()
    # tfidf words/col-names
    # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
```

```

for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model2.wv[word]
            # obtain the tf_idf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf

        except:
            pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    test_vectors.append(sent_vec)
    row += 1

```

In [0]: *#Standardization*

```

from sklearn.preprocessing import StandardScaler
sc= StandardScaler(with_mean=False)
X_train4 = sc.fit_transform(train_vectors)
X_test4 = sc.transform(test_vectors)

```

In [0]: *# creating odd list of K for KNN*

```

myList = list(range(0,50))
neighbors = list(filter(lambda x: x % 2 != 0, myList))

# empty list that will hold cv scores
cv_scores = []
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train4)]
# perform 10-fold cross validation
for k in neighbors:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train4, y_train, cv = my_cv, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_k = neighbors[MSE.index(min(MSE))]
print('\nThe optimal number of neighbors is %d.' % optimal_k)

```

```

# plot misclassification error vs k
plt.plot(neighbors, MSE)

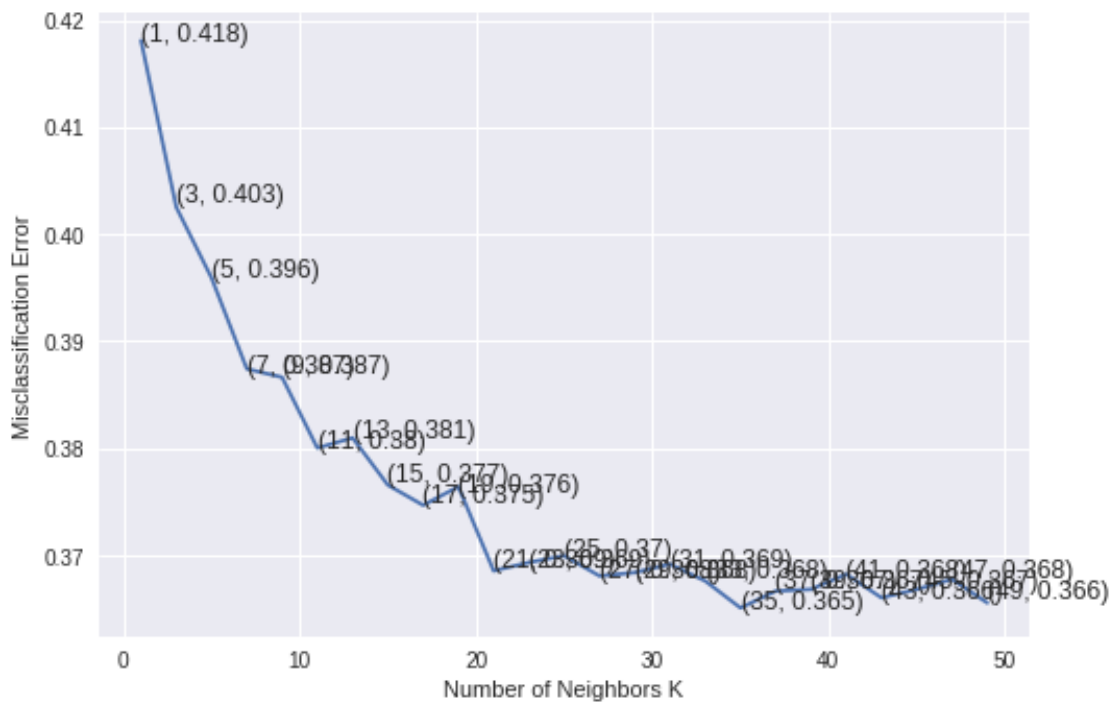
for xy in zip(neighbors, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('Misclassification Error')
plt.show()

print("the misclassification error for each k value is : ", (np.round(MSE,3)))

```

The optimal number of neighbors is 35.



```

the misclassification error for each k value is : [0.418 0.403 0.396 0.387 0.387 0.38 0.381 0.37
0.368 0.368 0.369 0.368 0.365 0.367 0.367 0.368 0.366 0.367 0.368
0.366]

```

### 6.0.1 Best 'k' = 35

```

In [76]: k = 35
         knn = KNeighborsClassifier(n_neighbors = k)

```

```

knn.fit(X_train4,y_train)
pred = knn.predict(X_test4)
acc = accuracy_score(y_test, pred, normalize=True) * float(100)
x = knn.predict(X_train4)
tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)
print('\n****Train accuracy for k = {} is {:.2f}'.format(k,tr_acc))
print('\n****Test accuracy for k = {} is {:.2f}'.format(k,acc))

```

```
****Train accuracy for k = 35 is 71.41
```

```
****Test accuracy for k = 35 is 47.37
```

## 6.1 Generate the confusion matrix

```
In [0]: print(confusion_matrix(y_test, pred))
```

```
[[ 105 1503]
 [   76 1316]]
```

```
In [0]: import plotly.tools as tls
```

```
tls.set_credentials_file(username='kv007', api_key='MSfBc0Jqh4Vq99anlb0p')
```

## 7 Conclusion

```
In [ ]: trace = go.Table(
```

```

    header=dict(values=['Algorithm', 'Model', 'Best Hyperparameter(K)', 'Train metric',
    cells=dict(values=[['Bag of words', 'Tf-idf', 'Average Word2Vec', 'Tf-idf Word2Vec',
                        ['K-NN', 'K-NN', 'K-NN', 'K-NN'],
                        [21,19,31,35],
                        ['accuracy(68.73%)', 'accuracy(48.40%)', 'accuracy(77.46%)', 'accuracy(71.41%)'],
                        ['accuracy(60.70%)', 'accuracy(46.40%)', 'accuracy(46.47%)', 'accuracy(47.37%)']])])

```

```
data = [trace]
```

```
py.iplot(data, filename = 'Conclusion')
```

Algorithm	Model	Best Hyperparameter(K)	Train metric	Test metric
Bag of words	K-NN	21	accuracy(68.73%)	accuracy(60.70%)
Tf-idf	K-NN	19	accuracy(48.40%)	accuracy(46.40%)
Average Word2Vec	K-NN	31	accuracy(77.46%)	accuracy(46.47%)
Tf-idf Word2Vec	K-NN	35	accuracy(71.41%)	accuracy(47.37%)